



# CS 554

# Computer Vision

---

## Homography & Image Stitching

**Hamdi Dibeklioglu**

Slide Credits: L. van der Maaten

# Image Stitching

- We are given a bunch of photographs; how do we *stitch* them together?



- But first, we need to understand *non-linear least squares* and *homographies*


# Non-linear least squares problems

# Non-linear least squares

- Fitting a panography was *linear* in the transformation parameters:

$$\begin{aligned} E &= \sum_i \|f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i\|^2 = \sum_i \|\mathbf{x}_i + J(\mathbf{x}_i)\mathbf{p} - \mathbf{x}'_i\|^2 \\ &= \sum_i \|J(\mathbf{x}_i)\mathbf{p} - \Delta\mathbf{x}_i\|^2 \end{aligned}$$

- For more complex motion models, the transformation is *non-linear*.






# Non-linear least squares

- Consider the *non-linear least squares problem*:

$$g(\mathbf{x}) = \|f(\mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2$$

- This problem is in general not *convex*; it may have multiple *local minima*:




# Taylor Expansion

- The *Taylor expansion* of the function  $f(x)$  around  $a$  is given by:

$$\sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$$

- Herein,  $f^{(n)}$  denotes the  $n$ -th derivative




# Gauss-Newton Method


- Iteratively find parameter updates  $\Delta \mathbf{x}$

# Gauss-Newton Method

- Iteratively find parameter updates  $\Delta \mathbf{x}$
- Perform a *first-order Taylor expansion* of the residual around the current  $\mathbf{x}$

$$\|f(\mathbf{x} - \Delta \mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2 \approx \|f(\mathbf{x}; \mathbf{A}) + J(\mathbf{x})\Delta \mathbf{x} - \mathbf{b}\|^2$$

  
parameter update  
objective

  
approximation around  $\mathbf{x}$


  
Jacobian




# Gauss-Newton Method

- Iteratively find parameter updates  $\Delta \mathbf{x}$
- Perform a *first-order Taylor expansion* of the residual around the current  $\mathbf{x}$

$$\|f(\mathbf{x} - \Delta \mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2 \approx \|f(\mathbf{x}; \mathbf{A}) + J(\mathbf{x})\Delta \mathbf{x} - \mathbf{b}\|^2$$

  
parameter update  
objective

  
approximation around  $\mathbf{x}$

  
Jacobian

- Note that the resulting residual approximation is linear in  $\Delta \mathbf{x}$  :
  - The parameter update  $\Delta \mathbf{x}$  may be obtained via linear least squares

# Gauss-Newton Method

- Writing down the linear least-squares solution for  $\Delta \mathbf{x}$ , we obtain:

$$\Delta \mathbf{x} = \left( J(\mathbf{x})^\top J(\mathbf{x}) \right)^{-1} J(\mathbf{x})^\top r(\mathbf{x})$$



“Gauss-Newton approximation to Hessian”

- Gauss-Newton iteratively performs this update:  $\mathbf{x} \leftarrow \mathbf{x} - \Delta \mathbf{x}$
- The Taylor expansion just became inaccurate! So iterate the whole process...

# Gauss-Newton Method

- Writing down the linear least-squares solution for  $\Delta \mathbf{x}$ , we obtain:

$$\Delta \mathbf{x} = \left( J(\mathbf{x})^\top J(\mathbf{x}) \right)^{-1} J(\mathbf{x})^\top r(\mathbf{x})$$



“Gauss-Newton approximation to Hessian”

- Gauss-Newton iteratively performs this update:  $\mathbf{x} \leftarrow \mathbf{x} - \Delta \mathbf{x}$
- The Taylor expansion just became inaccurate! So iterate the whole process...

- To implement, you only need to derive *Jacobian*:  $J(\mathbf{x}) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$

# Newton's Method

- Perform a *second-order Taylor expansion* of  $g(\mathbf{x})$  around  $\mathbf{x}$ :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^2] (\Delta\mathbf{x})^2$$

with residuals:  $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

# Newton's Method

- Perform a *second-order Taylor expansion* of  $g(\mathbf{x})$  around  $\mathbf{x}$ :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^2] (\Delta\mathbf{x})^2$$

with residuals:  $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

- This looks a lot like a linear least-squares problem; set gradient to zero:

$$-2J(\mathbf{x})^T r(\mathbf{x}) + 2 [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = 0$$

$$[H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = J(\mathbf{x})^T r(\mathbf{x})$$

# Newton's Method

- Perform a *second-order Taylor expansion* of  $g(\mathbf{x})$  around  $\mathbf{x}$ :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] (\Delta\mathbf{x})^2$$

with residuals:  $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

- This looks a lot like a linear least-squares problem; set gradient to zero:

$$-2J(\mathbf{x})^T r(\mathbf{x}) + 2 [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = 0$$


$$[H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = J(\mathbf{x})^T r(\mathbf{x})$$

- Note the similarity of the Newton update with the Gauss-Newton update:

$$\Delta\mathbf{x} = [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})]^{-1} J(\mathbf{x})^T r(\mathbf{x})$$

# Homography

# Let's design a camera




Idea 1: put a piece of film in front of an object

Do we get a reasonable image?




# Let's design a camera



Add a barrier to block off most of the rays


- This reduces blurring
- The opening is known as the **aperture**

# Pinhole camera



$f$  = focal length  
 $c$  = center of projection


# Modeling projection



The coordinate system

- We will use the pin-hole model as an approximation
- Put the optical center (**C**enter of projection) at the origin
- Where would you put the image (Projection Plane)?


# Modeling projection



## The coordinate system


- We will use the pin-hole model as an approximation
- Put the optical center (**C**enter of projection) at the origin
- Where would you put the image (**P**rojection **P**lane)?
- In front of the camera

# Modeling projection



## Projection equations

- Compute intersection with PP of ray from  $(x, y, z)$  to  $C$  in terms of  $x, y, z$ , and  $f$
- Hint: use similar triangles;  $ST = PQ * (TR/QR)$



$$(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z}, f)$$

- How do you get the 2D projection?

$$(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z})$$

Throw out the  $z$  coordinate

# Modeling projection

- *Focal length* of the camera influences what is captured on the image plane:



$$(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z}, f)$$

- A large focal length implies small *field of view*, and vice versa

# Homogeneous coordinates

- Is this a linear transformation?  $(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z})$

No; division by  $z$  is non-linear

Why?

Definition of a linear function:

$$f(ax) = a f(x)$$

$$f(x + y) = f(x) + f(y)$$

Lets look at a small example for  $f = 1/z$ :

$$1+2=3 \rightarrow f(1) + f(2) = f(1+2) \rightarrow \frac{1}{1} + \frac{2}{2} = \frac{3}{3} \rightarrow 1 + 1 = 1 \text{ ???}$$


Trick: add one more coordinate to remember what  $1/z$  was.

$$1+2=3 \rightarrow (\frac{1}{1}, 1) + (\frac{2}{2}, \frac{1}{2}) = (\frac{3}{3}, \frac{1}{3})$$

How would you use this coordinate to fix the problem?

Divide by the extra coordinate:  $1 + 2 = 3$

What does  $1/z$  look like?



# Homogeneous coordinates

- Is this a linear transformation?  $(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z})$

No; division by  $z$  is non-linear

Trick: add one more coordinate:

**2D:**

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image  
coordinates

**3D:**

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

homogeneous scene  
coordinates

Converting *from* homogeneous coordinates

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{w} \end{bmatrix} \Rightarrow (\tilde{x}/\tilde{w}, \tilde{y}/\tilde{w})$$

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{z} \\ \tilde{w} \end{bmatrix} \Rightarrow (\tilde{x}/\tilde{w}, \tilde{y}/\tilde{w}, \tilde{z}/\tilde{w})$$



# Homogeneous coordinates

- We are used to describing a location in *Cartesian coordinates*:

$$\mathbf{x} = [x \ y]^T \qquad \mathbf{x} = [x \ y \ z]^T$$

- Alternatively, we can describe locations in *homogeneous coordinates*:

$$\tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{w}]^T \qquad \tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{z} \ \tilde{w}]^T$$

- The corresponding Cartesian coordinates are given by:

$$\mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w}]^T \qquad \mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w} \ \tilde{z}/\tilde{w}]^T$$

- Essentially, you can think of  $\tilde{w}$  as a way to deal with object scale (“*disparity*”)
- Homogeneous coordinates are very useful when working with *perspective transformations* (*homographies*)

# Perspective Projection Matrix

Projection is a matrix multiplication using homogeneous coordinates:  $(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z})$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z/f \end{bmatrix} \Rightarrow (x \frac{f}{z}, y \frac{f}{z})$$

divide by the third  
coordinate

# Perspective Projection Matrix

Projection is a matrix multiplication using homogeneous coordinates:  $(x, y, z) \rightarrow (x \frac{f}{z}, y \frac{f}{z})$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z/f \end{bmatrix} \Rightarrow (x \frac{f}{z}, y \frac{f}{z})$$

In practice: split into lots of different coordinate transformations...

$$\begin{pmatrix} \text{2D} \\ \text{point} \\ (3 \times 1) \end{pmatrix} = \begin{pmatrix} \text{Camera to} \\ \text{pixel coord.} \\ \text{trans. matrix} \\ (3 \times 3) \end{pmatrix} \begin{pmatrix} \text{Perspective} \\ \text{projection matrix} \\ (3 \times 4) \end{pmatrix} \begin{pmatrix} \text{World to} \\ \text{camera coord.} \\ \text{trans. matrix} \\ (4 \times 4) \end{pmatrix} \begin{pmatrix} \text{3D} \\ \text{point} \\ (4 \times 1) \end{pmatrix}$$

# Camera Matrix

In practice: split into lots of different coordinate transformations...

$$\begin{pmatrix} \text{2D} \\ \text{point} \\ (3 \times 1) \end{pmatrix} = \begin{pmatrix} \text{Camera to} \\ \text{pixel coord.} \\ \text{trans. matrix} \\ (3 \times 3) \end{pmatrix} \begin{pmatrix} \text{Perspective} \\ \text{projection matrix} \\ (3 \times 4) \end{pmatrix} \begin{pmatrix} \text{World to} \\ \text{camera coord.} \\ \text{trans. matrix} \\ (4 \times 4) \end{pmatrix} \begin{pmatrix} \text{3D} \\ \text{point} \\ (4 \times 1) \end{pmatrix}$$

$$\tilde{\mathbf{x}} = \mathbf{K} [\mathbf{R} | \mathbf{t}] \mathbf{p} = \mathbf{P} \mathbf{p}$$


camera intrinsics  
(calibration matrix)

camera extrinsics  
(rotation + translation)

camera matrix


# How does a camera see a 3D point?

- The camera matrix is an example of a *projective transformation*:




# How does a camera see a 3D point?

What does this curve tell you about perspective effects?



# How does a camera see a 3D point?

What does this curve tell you about perspective effects?



Play limited role if the object is “far away”

# How does a camera see a place?

- Assume we have two cameras with projection matrices  $\tilde{\mathbf{P}}_0$  and  $\tilde{\mathbf{P}}_1$
- Where does camera 1 see the point that camera 0 sees at  $\tilde{\mathbf{x}}_0$ ?

$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{P}}_1 \tilde{\mathbf{P}}_0^{-1} \tilde{\mathbf{x}}_0 = \mathbf{M}_{10} \tilde{\mathbf{x}}_0$$



# How does a camera see a place?

- Assume we have two cameras with projection matrices  $\tilde{\mathbf{P}}_0$  and  $\tilde{\mathbf{P}}_1$
- Where does camera 1 see the point that camera 0 sees at  $\tilde{\mathbf{x}}_0$ ?

$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{P}}_1 \tilde{\mathbf{P}}_0^{-1} \tilde{\mathbf{x}}_0 = \mathbf{M}_{10} \tilde{\mathbf{x}}_0$$

- *Disparity* of point is irrelevant, so we can take the 3x3 sub-matrix of  $\mathbf{M}_{10}$ :


$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{H}}_{10} \tilde{\mathbf{x}}_0$$

- This is known as a *homography*  $\tilde{\mathbf{H}}_{10}$  between the two cameras

\* Notation in book is a bit sloppy here: page 56

# How does a camera see a place?


- Illustration of homography between two camera (for point and plane):



$$H_{10} = \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$$

# How does a camera see a place?

- Illustration of homography between two camera (for point and plane):



$$H_{10} = \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$$

- In Cartesian coordinates, the homography is given by:

$$x_1 = \frac{h_{00}x_0 + h_{01}y_0 + h_{02}}{h_{20}x_0 + h_{21}y_0 + 1}$$

$$y_1 = \frac{h_{10}x_0 + h_{11}y_0 + h_{12}}{h_{20}x_0 + h_{21}y_0 + 1}$$


# Image stitching

# Image stitching


- Stitching algorithms typically have four main ingredients:
  - Method to determine *correspondences* between images
  - Model describing the set of possible *motions* between images (homography)
  - Algorithm to perform *alignment* of the images (bundle adjustment)
  - Algorithm that *composites* the images after alignment (blending; seam finding)

# Determining correspondences


- Feature detection:



- Feature description:




- Feature matching:



# Motion model

- We will consider *homographies* because of their generality:



# Fitting a homography

- Recall the definition of a homography in Cartesian coordinates:

$$x' = f(x, y) = \frac{h_{00}x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad y' = g(x, y) = \frac{h_{10}x + h_{11}y + h_{12}}{h_{20}x + h_{21}y + 1}$$

- This makes the alignment objective a non-linear least squares problem:

$$\sum_i \left\| \begin{bmatrix} f(x_i, y_i; \mathbf{H}) \\ g(x_i, y_i; \mathbf{H}) \end{bmatrix} - \begin{bmatrix} x'_i \\ y'_i \end{bmatrix} \right\|^2$$



# Fitting a homography

- Recall the definition of a homography in Cartesian coordinates:

$$x' = f(x, y) = \frac{h_{00}x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad y' = g(x, y) = \frac{h_{10}x + h_{11}y + h_{12}}{h_{20}x + h_{21}y + 1}$$

- This makes the alignment objective a non-linear least squares problem:

$$\sum_i \left\| \begin{bmatrix} f(x_i, y_i; \mathbf{H}) \\ g(x_i, y_i; \mathbf{H}) \end{bmatrix} - \begin{bmatrix} x'_i \\ y'_i \end{bmatrix} \right\|^2$$

- Simply use Gauss-Newton's (or Newton's) method to solve this problem:
  - You may have to use RANSAC to deal with outliers!

# Remark on RANSAC

- The treatment of RANSAC we saw last week was somewhat simplified
- It was not random at all!

# Remark on RANSAC

- The treatment of RANSAC we saw last week was somewhat simplified
- It was not random at all! Full RANSAC actually works as follows:
  - 1) Select random subset from 50% of “best” matches as initial inliers
  - 2) Model is fitted to the *hypothetical inliers*
  - 3) Data are tested against the fitted model to determine hypothetical inliers
  - 4) Return to step 2) until sufficient points are classified as inliers  
(or fixed number of times)
  - 5) Return to step 1) a fixed number of times


# Fitting a homography

- Consider the following two images and SIFT matches between them:




# Fitting a homography

- Consider the following two images and SIFT matches between them:



# Fitting a homography

- Visualization of the inliers found by RANSAC:





# Fitting a homography

- Visualization of the homography found between the two images:



# Alpha blending


- *Alpha blending* averages the images; leads to “ghosting” and visible seams:





# Finding an optimal seam

- We would like to find a *seam* that minimizes the difference between images
- This can be formulated as a *graph min-cut* problem, as follows:



- When you want pixel  $\mathbf{x}$  to be taken from  $I_1$ , set  $f_2(\mathbf{x}) = \infty$  (and vice versa)
- In all other cases:  $f_1(\mathbf{x}) = f_2(\mathbf{x}) = 0$
- Set, e.g.,  $g(x, y) = |I_1(\mathbf{x}) - I_2(\mathbf{y})| + |I_2(\mathbf{x}) - I_1(\mathbf{y})|$
- Efficient graph-cut algorithms exist

# Finding an optimal seam

- Ghosting can be prevented by finding a *seam* at which to switch images:




# Feathering

- Visibility of seam may be reduced by “*blurring*” weights of both images:



# Correcting for exposure differences

- Instead of using seam to select image, we use it to select the *image gradient*.






# Correcting for exposure differences

- The stitched image can then be recovered from this gradient:



# Lens distortion

- Homography model may be too simple when lens distortion is present:



- A common way of dealing with this is by incorporating a distortion model in the objective, and also minimizing w.r.t. the parameters of that model

# Other applications of homographies




# Other applications of homographies







# Other applications: Image Warping



How would you do this?


Source: Steve Seitz

# Other applications: Image rectification



# Other applications: Analyzing patterns/shapes

What is the shape of the b/w floor pattern?



Homography




Slide: Criminisi

**The floor (enlarged)**

**Automatically rectified floor**

# Other applications: Analyzing patterns/shapes

Automatic rectification



From Martin Kemp *The Science of Art*  
(*manual reconstruction*)



# Other applications: Analyzing patterns/shapes



*St. Lucy Altarpiece, D. Veneziano*




**Automatically rectified floor**

# Other applications: Analyzing patterns/shapes



**Automatic  
rectification**



**From Martin Kemp, *The Science of Art*  
(*manual reconstruction*)**

Reading material: Section 2 and 9 of Szeliski