

# CS 554

# Computer Vision

---


## Tracking

**Hamdi Dibeklioglu**

Slide Credits: L. van der Maaten

# Tracking

Tracking aims to follow an object or object configuration in a video sequence:



# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize  $posterior = likelihood \times prior$ ?

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior = likelihood x prior*?

Recall *Bayes' rule*: 
$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \frac{p(\mathbf{I}_t | \mathbf{x}_t, \mathbf{x}_{t-1}) p(\mathbf{x}_t | \mathbf{x}_{t-1})}{p(\mathbf{I}_t | \mathbf{x}_{t-1})}$$


# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior = likelihood x prior*?

Recall *Bayes' rule*:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \frac{p(\mathbf{I}_t | \mathbf{x}_t, \mathbf{x}_{t-1}) p(\mathbf{x}_t | \mathbf{x}_{t-1})}{p(\mathbf{I}_t | \mathbf{x}_{t-1})}$$



location and size  
of object at time  $t$

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior = likelihood x prior*?

Recall *Bayes' rule*:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \frac{p(\mathbf{I}_t | \mathbf{x}_t, \mathbf{x}_{t-1}) p(\mathbf{x}_t | \mathbf{x}_{t-1})}{p(\mathbf{I}_t | \mathbf{x}_{t-1})}$$

location and size  
of object at time  $t$

observed image  
at time  $t$



# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior* = *likelihood*  $\times$  *prior*?

Recall *Bayes' rule*:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \frac{p(\mathbf{I}_t | \mathbf{x}_t, \mathbf{x}_{t-1}) p(\mathbf{x}_t | \mathbf{x}_{t-1})}{p(\mathbf{x}_{t-1})}$$

location and size of object at time  $t$

observed image at time  $t$

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior* = *likelihood*  $\times$  *prior*?

Recall *Bayes' rule*:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \frac{p(\mathbf{I}_t | \mathbf{x}_t) p(\mathbf{x}_t | \mathbf{x}_{t-1})}{p(\mathbf{I}_t)}$$

location and size of object at time  $t$


observed image at time  $t$

# Tracking

Every tracker consists of three main components:

- *Likelihood*: How well does a configuration explain the current observation?
- *Prior*: Given the previous object location, how likely is a configuration?
- *Search strategy*: How do we maximize *posterior* = *likelihood*  $\times$  *prior*?

Recall *Bayes' rule*:  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t | \mathbf{x}_t) p(\mathbf{x}_t | \mathbf{x}_{t-1})$



location and size  
of object at time  $t$

observed image  
at time  $t$

# Tracking

*Appearance likelihood* can be any object-recognition model

# Tracking

*Appearance likelihood* can be any object-recognition model

*Motion prior* can be a linear dynamical system

# Tracking

*Appearance likelihood* can be any object-recognition model

*Motion prior* can be a linear dynamical system

*Search strategy* can be one of the following four approaches:

- Kalman (only for linear-Gaussian likelihood and prior) or particle filtering
- Lucas-Kanade algorithm (only for template matching using squared errors)
- Mean-shift tracking
- Sliding-window search (essentially brute-force search)

Kalman filter

# Kalman filter

- We aim to infer the location of the object,  $\mathbf{x}_t$ , at time  $t$
- A simple *motion prior* could be defined as:

$$p(\mathbf{x}_t|\mathbf{x}_{t-1}) \propto \exp\left(-\frac{1}{2\sigma^2}\|\mathbf{x}_t - \mathbf{A}\mathbf{x}_{t-1}\|^2\right)$$

- A simple *likelihood* model could be defined as:

$$p(\mathbf{I}_t|\mathbf{x}_t) \propto \exp\left(-\frac{1}{2\tau^2}\|\mathbf{B}\mathbf{x}_t - \phi(\mathbf{I}_t)\|^2\right)$$

- Combining the two via Bayes' rule:  $p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t|\mathbf{x}_t)p(\mathbf{x}_t|\mathbf{x}_{t-1})$



# Kalman filter

- Combining the two via Bayes' rule:  $p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t|\mathbf{x}_t)p(\mathbf{x}_t|\mathbf{x}_{t-1})$
- Because both terms are Gaussian, the result is also Gaussian:

$$p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) = \mathcal{N}(\mathbf{x}_t|\boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)$$

# Kalman filter

- Combining the two via Bayes' rule:  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t | \mathbf{x}_t) p(\mathbf{x}_t | \mathbf{x}_{t-1})$
- Because both terms are Gaussian, the result is also Gaussian:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}) \propto \exp \left( -\frac{1}{2\sigma^2} \|\mathbf{x}_t - \mathbf{A}\mathbf{x}_{t-1}\|^2 \right)$$

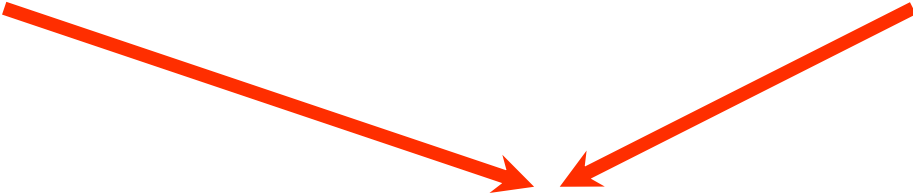


$$p(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t | \mathbf{A}\mathbf{x}_{t-1}, \sigma^2)$$

$$p(\mathbf{I}_t | \mathbf{x}_t) \propto \exp \left( -\frac{1}{2\tau^2} \|\mathbf{B}\mathbf{x}_t - \phi(\mathbf{I}_t)\|^2 \right)$$



$$p(\mathbf{I}_t | \mathbf{x}_t) = \mathcal{N}(\phi(\mathbf{I}_t) | \mathbf{B}\mathbf{x}_t, \tau^2)$$


$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) = \mathcal{N}(\mathbf{x}_t | \boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)$$

# Kalman filter

- Combining the two via Bayes' rule:  $p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t|\mathbf{x}_t)p(\mathbf{x}_t|\mathbf{x}_{t-1})$
- Because both terms are Gaussian, the result is also Gaussian:

$$p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) = \mathcal{N}(\mathbf{x}_t|\boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)$$

# Kalman filter

- Combining the two via Bayes' rule:  $p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) \propto p(\mathbf{I}_t|\mathbf{x}_t)p(\mathbf{x}_t|\mathbf{x}_{t-1})$
- Because both terms are Gaussian, the result is also Gaussian:

$$p(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{I}_t) = \mathcal{N}(\mathbf{x}_t|\boldsymbol{\mu}_t, \boldsymbol{\Sigma}_t)$$

$$\boldsymbol{\mu}_t = \mathbf{A}\boldsymbol{\mu}_{t-1} + \mathbf{K}_t(\phi(\mathbf{I}_t) - \mathbf{B}\mathbf{A}\boldsymbol{\mu}_{t-1})$$

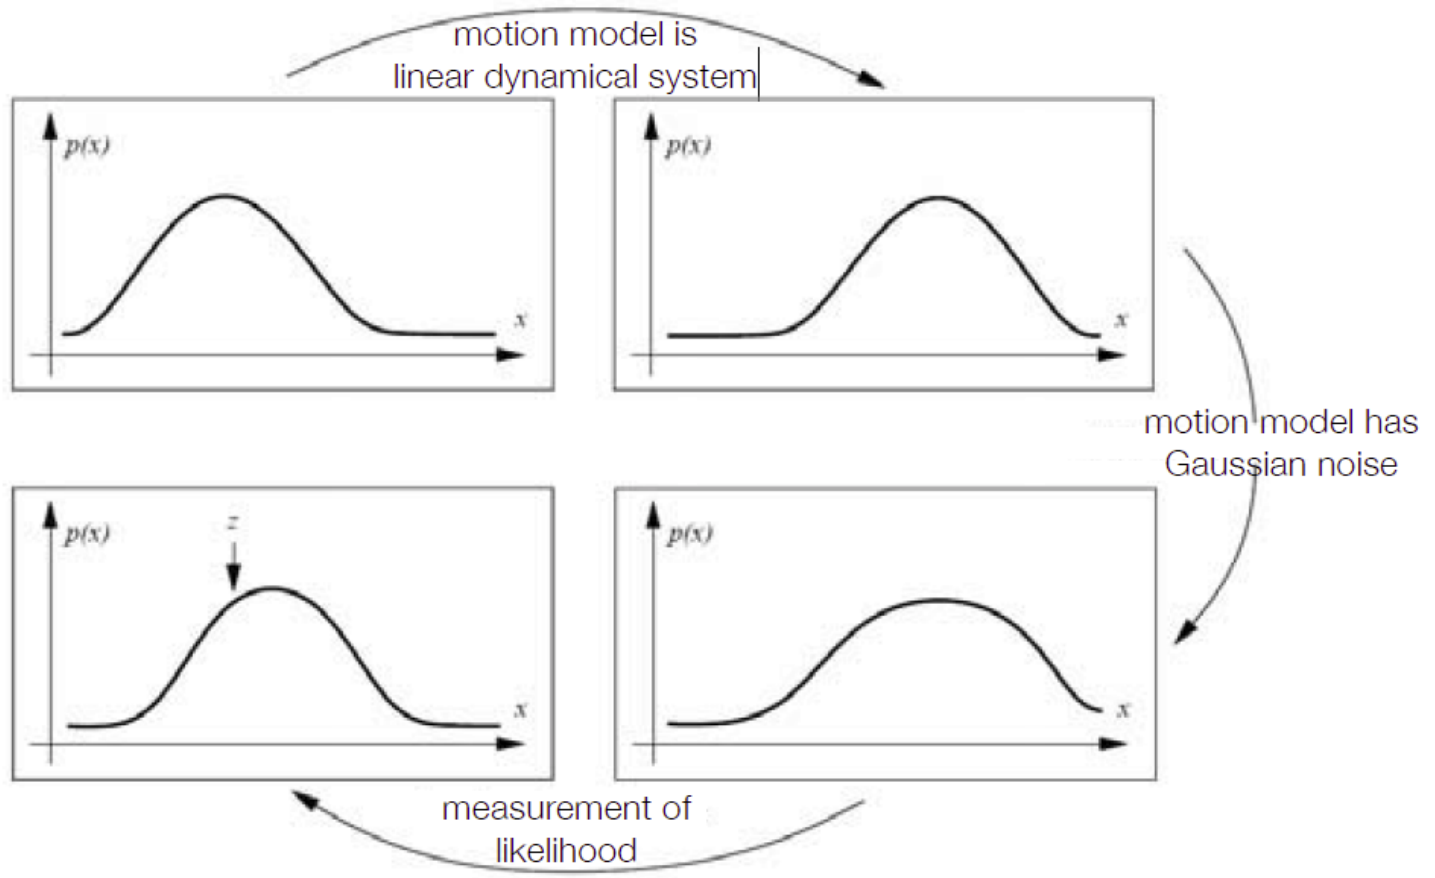
$$\boldsymbol{\Sigma}_t = (\mathbf{I} - \mathbf{K}_t\mathbf{B}) (\mathbf{A}\boldsymbol{\Sigma}_{t-1}\mathbf{A}^\top + \sigma^2\mathbf{I})$$

$$\mathbf{K}_t = (\mathbf{A}\boldsymbol{\Sigma}_{t-1}\mathbf{A}^\top + \sigma^2\mathbf{I}) \mathbf{B}^\top [\mathbf{B} (\mathbf{A}\boldsymbol{\Sigma}_{t-1}\mathbf{A}^\top + \sigma^2\mathbf{I}) \mathbf{B}^\top + \tau^2\mathbf{I}]^{-1}$$



"Kalman gain matrix"


# Kalman filter



# Example: Kalman filter



# Example: Kalman filter



# Kalman filter

In practical settings, the Kalman filter does often not work very well:

- The Gaussian likelihood model is generally way too simple in practice:  
Gaussians are unimodal and can therefore only maintain a single hypothesis



# Kalman filter

In practical settings, the Kalman filter does often not work very well:

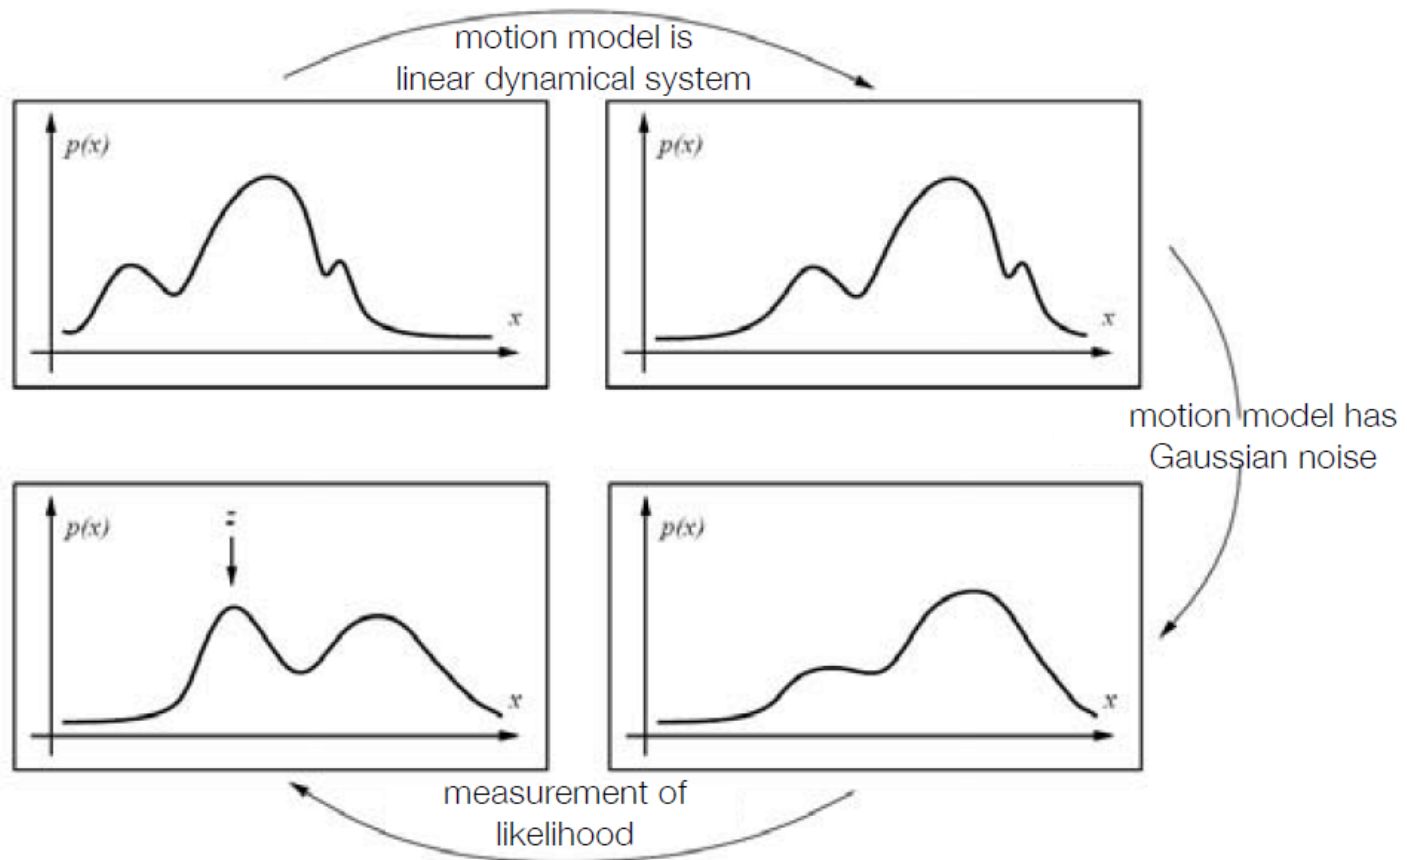
- The Gaussian likelihood model is generally way too simple in practice:  
Gaussians are unimodal and can therefore only maintain a single hypothesis

Changing the likelihood to something non-Gaussian complicates inference:

- Particle filters do this: solve inference problem by *importance sampling*


Particle filter

# Particle filter



# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters represent  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  by a set of weighted samples (“particles”):



# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:

# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:
  - Assume we have a set of  $N$  weighted samples  $\{(\mathbf{s}_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$

# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:
  - Assume we have a set of  $N$  weighted samples  $\{(\mathbf{s}_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$
  - Pick sample  $\mathbf{s}_{t-1}^{(n)}$  according to weights, and sample from  $p(\mathbf{s}_t^{(n)} | \mathbf{s}_{t-1}^{(n)})$  (repeat  $N$  times)

# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:
  - Assume we have a set of  $N$  weighted samples  $\{(\mathbf{s}_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$
  - Pick sample  $\mathbf{s}_{t-1}^{(n)}$  according to weights, and sample from  $p(\mathbf{s}_t^{(n)} | \mathbf{s}_{t-1}^{(n)})$  (repeat  $N$  times)
  - Reweight samples according to  $\pi_t^{(n)} \leftarrow \pi_{t-1}^{(n)} p(\mathbf{I}_t | \mathbf{s}_t^{(n)})$



# Particle filter (Condensation)

- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:
  - Assume we have a set of  $N$  weighted samples  $\{(\mathbf{s}_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$
  - Pick sample  $\mathbf{s}_{t-1}^{(n)}$  according to weights, and sample from  $p(\mathbf{s}_t^{(n)} | \mathbf{s}_{t-1}^{(n)})$  (repeat  $N$  times)
  - Reweight samples according to  $\pi_t^{(n)} \leftarrow \pi_{t-1}^{(n)} p(\mathbf{I}_t | \mathbf{s}_t^{(n)})$
  - Renormalize samples by  $\pi_t^{(n)} \leftarrow \frac{\pi_t^{(n)}}{\sum_{n=1}^N \pi_t^{(n)}}$

# Particle filter (Condensation)


- For complex appearance likelihood functions, exact inferences are impossible
- Particle filters sample from  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t)$  via *sequential importance sampling*:
  - Assume we have a set of  $N$  weighted samples  $\{(\mathbf{s}_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$
  - Pick sample  $\mathbf{s}_{t-1}^{(n)}$  according to weights, and sample from  $p(\mathbf{s}_t^{(n)} | \mathbf{s}_{t-1}^{(n)})$  (repeat  $N$  times)
  - Reweight samples according to  $\pi_t^{(n)} \leftarrow \pi_{t-1}^{(n)} p(\mathbf{I}_t | \mathbf{s}_t^{(n)})$
  - Renormalize samples by  $\pi_t^{(n)} \leftarrow \frac{\pi_t^{(n)}}{\sum_{n=1}^N \pi_t^{(n)}}$
  - Prediction may be weighted average of particles, or most likely particle

# Particle filter (Condensation)

Pick sample  $\mathbf{s}_{t-1}^{(n)}$   
according to weights

Sample from  $p(\mathbf{s}_t^{(n)} | \mathbf{s}_{t-1}^{(n)})$

Update weights  $\pi_t^{(n)}$



# Example: Particle filter



Tracking-by-detection

# Tracking-by-detection

- Appearance likelihood: Your favorite object detector (V&J, D-T, *etc.*)
  - Now, we are using a *conditional likelihood*:

$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) \propto \exp \{ s_L(\mathbf{x}_t, \mathbf{I}_t) + s_M(\mathbf{x}_t, \mathbf{x}_{t-1}) \}$$

- Motion prior: May be virtually any motion model
- Search strategy: Sliding window detector at multiple scales (brute-force)

# Tracking-by-detection

- Appearance likelihood: Your favorite object detector (V&J, D-T, *etc.*)
  - Now, we are using a *conditional likelihood*:


$$p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{I}_t) \propto \exp \{ s_L(\mathbf{x}_t, \mathbf{I}_t) + s_M(\mathbf{x}_t, \mathbf{x}_{t-1}) \}$$

- Motion prior: May be virtually any motion model
- Search strategy: Sliding window detector at multiple scales (brute-force)
- Potential problem: Object appearance may change over time


# Track-learn-detect

Can we gather positive and negative examples to update appearance model?

- Positive example: Assume the track in the previous frame is correct
- Negative example: Example with high detection score but low motion prior



appearance likelihood




motion prior




# Track-learn-detect

Can we gather positive and negative examples to update appearance model?


- Positive example: Assume the track in the previous frame is correct
- Negative example: Example with high detection score but low motion prior



appearance likelihood



motion prior




posterior over object location


# Track-learn-detect

Can we gather positive and negative examples to update appearance model?


- Positive example: Assume the track in the previous frame is correct
- Negative example: Example with high detection score but low motion prior



appearance likelihood



motion prior




posterior over object location


# Track-learn-detect

Can we gather positive and negative examples to update appearance model?


- Positive example: Assume the track in the previous frame is correct
- Negative example: Example with high detection score but low motion prior



appearance likelihood



motion prior



posterior over object location

# Example: Track-learn-detect



Kanade-Lucas-Tomasi tracker

# Kanade-Lucas-Tomasi Tracker

- Detect feature points in object using *Shi-Tomasi corner detector* (like Harris)
- Track feature points in the next frame by minimizing the squared error:

$$\sum_{n=1}^N \left[ \mathbf{I}_{t-1} \left( \mathbf{x}_{t-1}^{(n)} \right) - \mathbf{I}_t \left( W \left( \mathbf{x}_{t-1}^{(n)}; \mathbf{p} \right) \right) \right]^2$$

# Kanade-Lucas-Tomasi Tracker

- Detect feature points in object using *Shi-Tomasi corner detector* (like Harris)
- Track feature points in the next frame by minimizing the squared error:

$$\sum_{n=1}^N \left[ \mathbf{I}_{t-1} \left( \mathbf{x}_{t-1}^{(n)} \right) - \mathbf{I}_t \left( W \left( \mathbf{x}_{t-1}^{(n)}; \mathbf{p} \right) \right) \right]^2$$

  
small image patch around  
feature point at time  $t-1$

# Kanade-Lucas-Tomasi Tracker

- Detect feature points in object using *Shi-Tomasi corner detector* (like Harris)
- Track feature points in the next frame by minimizing the squared error:

$$\sum_{n=1}^N \left[ \mathbf{I}_{t-1} \left( \mathbf{x}_{t-1}^{(n)} \right) - \mathbf{I}_t \left( W \left( \mathbf{x}_{t-1}^{(n)}; \mathbf{p} \right) \right) \right]^2$$

small image patch around  
feature point at time  $t-1$

small image patch around  
warped feature point at time  $t$



# Kanade-Lucas-Tomasi Tracker

- Detect feature points in object using *Shi-Tomasi corner detector* (like Harris)
- Track feature points in the next frame by minimizing the squared error:

$$\sum_{n=1}^N \left[ \mathbf{I}_{t-1} \left( \mathbf{x}_{t-1}^{(n)} \right) - \mathbf{I}_t \left( W \left( \mathbf{x}_{t-1}^{(n)}; \mathbf{p} \right) \right) \right]^2$$

small image patch around  
feature point at time  $t-1$

small image patch around  
warped feature point at time  $t$

- Note that this is yet another (non)linear least squares minimization problem!

# Kanade-Lucas-Tomasi Tracker

- Detect feature points in object using *Shi-Tomasi corner detector* (like Harris)
- Track feature points in the next frame by minimizing the squared error:

$$\sum_{n=1}^N \left[ \mathbf{I}_{t-1} \left( \mathbf{x}_{t-1}^{(n)} \right) - \mathbf{I}_t \left( W \left( \mathbf{x}_{t-1}^{(n)}; \mathbf{p} \right) \right) \right]^2$$

small image patch around  
feature point at time  $t-1$

small image patch around  
warped feature point at time  $t$

- Note that this is yet another (non)linear least squares minimization problem!
  - The minimization can be performing using the Lucas-Kanade algorithm

# Example: Kanade-Lucas-Tomasi Tracker



## **Reading material:**

- M. Isard and A. Blake. "Condensation—Conditional Density Propagation for Visual Tracking." International Journal of Computer Vision 29(1), 5-28, 1998.
- Section 1 and 2 of "Lucas-Kanade 20 Years On: A Unifying Framework," International Journal of Computer Vision 56(3), 221–255, 2004.